

```
In [1]: #importy
from pathlib import Path
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans

from sklearn.preprocessing import StandardScaler
import joblib
```

```
In [2]: #cesta
CSV_PATH = Path(r"C:\Users\adria\PycharmProjects\pythonProject\ZNEUS\PROJEKT\houses
CSV_PATH
```

```
Out[2]: WindowsPath('C:/Users/adria/PycharmProjects/pythonProject/ZNEUS/PROJEKT/houses.csv')
```

EDA - Data analysis

Load dataset

```
In [3]: df = pd.read_csv(CSV_PATH)

print("Shape:", df.shape)
display(df.head(10))
```

Shape: (20640, 9)

	median_house_value	median_income	housing_median_age	total_rooms	total_bedrooms	...
0	452600.0	8.3252	41.0	880.0	129.0	
1	358500.0	8.3014	21.0	7099.0	1106.0	
2	352100.0	7.2574	52.0	1467.0	190.0	
3	341300.0	5.6431	52.0	1274.0	235.0	
4	342200.0	3.8462	52.0	1627.0	280.0	
5	269700.0	4.0368	52.0	919.0	213.0	
6	299200.0	3.6591	52.0	2535.0	489.0	
7	241400.0	3.1200	52.0	3104.0	687.0	
8	226700.0	2.0804	42.0	2555.0	665.0	
9	261100.0	3.6912	52.0	3549.0	707.0	

Dtypes, missing, duplicates

```
In [4]: display(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   median_house_value    20640 non-null   float64
 1   median_income        20640 non-null   float64
 2   housing_median_age   20640 non-null   float64
 3   total_rooms          20640 non-null   float64
 4   total_bedrooms       20640 non-null   float64
 5   population          20640 non-null   float64
 6   households          20640 non-null   float64
 7   latitude             20640 non-null   float64
 8   longitude            20640 non-null   float64
dtypes: float64(9)
memory usage: 1.4 MB
None
```

```
In [5]: print("Dtypes, missing counts")
dtypes_missing = pd.DataFrame({
    "dtype": df.dtypes.astype(str),
    "missing_count": df.isna().sum(),
    "missing_rate": (df.isna().mean()).round(4),
    "unique_count": df.nunique(dropna=True)
}).sort_values("missing_count", ascending=False)
display(dtypes_missing)

dups = df.duplicated().sum()
print("Duplicated rows:", dups)
```

Dtypes, missing counts

		dtype	missing_count	missing_rate	unique_count
	median_house_value	float64	0	0.0	3842
	median_income	float64	0	0.0	12928
	housing_median_age	float64	0	0.0	52
	total_rooms	float64	0	0.0	5926
	total_bedrooms	float64	0	0.0	1928
	population	float64	0	0.0	3888
	households	float64	0	0.0	1815
	latitude	float64	0	0.0	862
	longitude	float64	0	0.0	844

Duplicated rows: 0

Statistics for numeric columns

```
In [6]: num = df.select_dtypes(include=[np.number])
if num.shape[1] > 0:
    display(num.describe().T)
else:
    print("No numeric columns found.")
```

	count	mean	std	min	25%
median_house_value	20640.0	206855.816909	115395.615874	14999.0000	119600.0000
median_income	20640.0	3.870671	1.899822	0.4999	2.5634
housing_median_age	20640.0	28.639486	12.585558	1.0000	18.0000
total_rooms	20640.0	2635.763081	2181.615252	2.0000	1447.7500
total_bedrooms	20640.0	537.898014	421.247906	1.0000	295.0000
population	20640.0	1425.476744	1132.462122	3.0000	787.0000
households	20640.0	499.539680	382.329753	1.0000	280.0000
latitude	20640.0	35.631861	2.135952	32.5400	33.9300
longitude	20640.0	-119.569704	2.003532	-124.3500	-121.8000

```
In [7]: maxv = df['median_house_value'].max()
print("Max value:", maxv, " Count:", (df['median_house_value']==maxv).sum())
df['is_censored'] = (df['median_house_value']==maxv).astype(int)
df['is_censored'].value_counts()
```

Max value: 500001.0 Count: 965

```
Out[7]: is_censored
0    19675
1      965
Name: count, dtype: int64
```


```
In [8]: if num.shape[1] > 0:
    iqr = num.quantile(0.75) - num.quantile(0.25)
    range_ = num.max() - num.min()
    outlier_flag = (range_ / (iqr.replace(0, np.nan))).sort_values(ascending=False)
    print("\nTop features by range / IQR (may indicate heavy tails or outliers):")
    display(outlier_flag)
else:
    print("No numeric columns found.")
```

Top features by range / IQR (may indicate heavy tails or outliers):

```
population           38.037313
total_rooms          23.124835
households           18.710769
total_bedrooms       18.306818
median_income        6.651926
median_house_value   3.341960
housing_median_age   2.684211
longitude            2.649077
latitude             2.489418
dtype: float64
```

Correlation matrix for numeric features

```
In [9]: num = df.select_dtypes(include=[np.number])
if num.shape[1] >= 2:
    corr = num.corr(numeric_only=True)
    plt.figure(figsize=(8,6))
    plt.imshow(corr, interpolation="nearest", aspect="auto")
    plt.colorbar()
    plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
    plt.yticks(range(len(corr.index)), corr.index)
    plt.title("Correlation matrix (numeric)")
    plt.tight_layout()
    plt.show()
else:
    print("Not enough numeric columns for correlation matrix.")
```



```
In [10]: target = "median_house_value"
if target in corr.columns:
    target_corr = corr[target].drop(target).abs().sort_values(ascending=False)
    print("Top features by absolute correlation with target:")
    display(target_corr.head(10))
else:
    print("Target not in numeric columns for correlation.")
```

Top features by absolute correlation with target:

median_income	0.688075
is_censored	0.562614
latitude	0.144160
total_rooms	0.134153
housing_median_age	0.105623
households	0.065843
total_bedrooms	0.050594
longitude	0.045967
population	0.024650



Name: median_house_value, dtype: float64

Target distribution

```
In [11]: plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
df[target].dropna().hist(bins=40)
plt.title(f"Histogram: {target}")
plt.xlabel(target)

plt.subplot(1,2,2)
df[target].dropna().plot.box(vert=False)
plt.title(f"Boxplot: {target}")

plt.tight_layout()
plt.show()
```



```
In [12]: skew = num.skew().sort_values(ascending=False)
display(skew)
```

```

population           4.935858
is_censored         4.294220
total_rooms          4.147343
total_bedrooms       3.453073
households           3.410438
median_income        1.646657
median_house_value   0.977763
latitude              0.465953
housing_median_age   0.060331
longitude             -0.297801
dtype: float64

```

```
In [13]: kurtosis = num.kurtosis().sort_values(ascending=False)
display(kurtosis)
```

```


population           73.553116
total_rooms          32.630927
households           22.057988
total_bedrooms       21.923495
is_censored          16.441921
median_income         4.952524
median_house_value   0.327870
housing_median_age   -0.800629
latitude              -1.117760
longitude             -1.330152
dtype: float64

```

Scatter and Pair plots for top numeric features correlated with target


```
In [14]: if target in num.columns:
    corr_with_target = num.corr()[target].abs().drop(target).sort_values(ascending=True)
    top_feats = corr_with_target.head(3).index.tolist()
    print("Top numeric features vs target:", top_feats)
    n = len(top_feats)
    if n > 0:
        plt.figure(figsize=(5*n,3))
        for i, f in enumerate(top_feats,1):
            plt.subplot(1,n,i)
            plt.scatter(df[f], df[target], s=8, alpha=0.6)
            plt.xlabel(f); plt.ylabel(target)
            plt.title(f"{f} vs {target}")
        plt.tight_layout()
        plt.show()
    else:
        print("Target not numeric or no numeric features.")
```

Top numeric features vs target: ['median_income', 'is_censored', 'latitude']



```
In [15]: sns.pairplot(df.sample(2000)[top_feats])
```

Out[15]: <seaborn.axisgrid.PairGrid at 0x170b4a4b7a0>




Region exploration

```
In [16]: coords = df[['latitude','longitude']]
km = KMeans(n_clusters=8, random_state=0).fit(coords)
df['region_id'] = km.labels_
df.groupby('region_id')['median_house_value'].median().sort_values()
```

```
Out[16]: region_id
2    76400.0
4    85800.0
6    89000.0
3   117800.0
7   169650.0
5   181600.0
0   198400.0
1   245100.0
Name: median_house_value, dtype: float64
```

```
In [17]: plt.figure(figsize=(8,6))
plt.scatter(df['longitude'], df['latitude'], c=df['region_id'], s=8, cmap='tab10',
plt.gca().invert_xaxis()
plt.title("KMeans clusters (region_id)")
plt.xlabel("longitude"); plt.ylabel("latitude")
plt.show()
```




```
In [18]: df.groupby('region_id', observed=False)[ 'median_house_value' ].agg([ 'count', 'median'
```

Out[18]:

	count	median	mean
region_id			
2	1117	76400.0	86213.249776
4	685	85800.0	92102.916788
6	474	89000.0	106898.529536
3	2203	117800.0	127250.976396
7	1678	169650.0	198141.446365
5	986	181600.0	191438.676471
0	8652	198400.0	229496.027624
1	4845	245100.0	262594.436739

In [19]:

```
centroids = km.cluster_centers_
plt.figure(figsize=(8,6))
plt.hexbin(df["longitude"], df["latitude"], gridsize=80, mincnt=1, cmap='Greys', alpha=0.5)
plt.scatter(centroids[:,1], centroids[:,0], c='red', s=100, marker='X')
for i,(lon,lat) in enumerate(zip(centroids[:,1], centroids[:,0])):
    plt.text(lon, lat, str(i), color='white', fontsize=9, ha='center', va='center')
plt.gca().invert_xaxis()
plt.show()
```



Categorical check

```
In [20]: categ = df.select_dtypes(include=["object","category"]).columns.tolist()
if not categ:
    print("No categorical columns found.")
else:
    for c in categ:
        print(f"\nColumn: {c} | Unique: {df[c].nunique()}")
        print(df[c].value_counts().head(5))
```

No categorical columns found.

EDA zhrnutie:

- **Veľkosť a čistota dát:** 20 640 riadkov, 9 stĺpcov. V EDA žiadne chýbajúce hodnoty a 0 úplných duplicitných riadkov — dáta sú kompletné a pripravené na ďalšiu analýzu.
- **Ciel — median_house_value :**
 - mean ≈ **206 856**, median ≈ **179 700**, veľký rozptyl.
 - max = **500001.0**, počet pri maxime = **965**
- **Hlavné prediktory (lineárne):**
 - `median_income` má najsilnejšiu pozitívnu koreláciu s cenou (~**0.69**) — kľúčový prediktor.
 - Ďalej sú relevantné priestorové prvky (`latitude` / `longitude`) a štrukturálne metriky domu `total_rooms` ...
- **Outliery a šikmost:**
 - Stĺpce `population`, `total_rooms`, `households`, `total_bedrooms` majú veľký `range` / `IQR` a vysoký skew → silné chvosty.
 - `median_house_value` skew ≈ **0.98** (pravostranný).
- **Vizualizácie a priestorový vzor:**
 - Scatter / pairplot potvrdili silný trend `median_income` vs cena a regionálne pásy podľa `latitude`.
 - Hexbin (`longitude` vs `latitude`) ukázal jasné **hotspots**
- **Regionálne klastre (KMeans, k=8):**
 - KMeans vytvoril regiony s rôznymi mediánmi
 - Veľkosti klastrov sú nerovnomerné (niektoré stovky, iné tisícky riadkov) — pri použití v modeli zohľadni nerovnomernosť.
 - Dôležité: KMeans v EDA možno fitnúť na celom df len pre exploráciu

Data preprocessing and normalization

