

```
In [1]: #importy
from pathlib import Path
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.model_selection import train_test_split
import joblib
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import FunctionTransformer, StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.dummy import DummyRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

import os, json, uuid, time, math, random
from pathlib import Path

import torch
from torch import nn
from torch.utils.data import TensorDataset, DataLoader
from torch.amp import autocast, GradScaler
scaler = GradScaler('cuda')
```

```
In [2]: #cesta
CSV_PATH = Path(r"C:\Users\adria\PycharmProjects\pythonProject\ZNEUS\PROJEKT\houses
CSV_PATH
```

```
Out[2]: WindowsPath('C:/Users/adria/PycharmProjects/pythonProject/ZNEUS/PROJEKT/houses.cs
v')
```

EDA - Data analysis

Load dataset

```
In [3]: df = pd.read_csv(CSV_PATH)

print("Shape:", df.shape)
display(df.head(10))
```

```
Shape: (20640, 9)
```

| | median_house_value | median_income | housing_median_age | total_rooms | total_bedrooms | |
|---|--------------------|---------------|--------------------|-------------|----------------|--|
| 0 | 452600.0 | 8.3252 | 41.0 | 880.0 | 129.0 | |
| 1 | 358500.0 | 8.3014 | 21.0 | 7099.0 | 1106.0 | |
| 2 | 352100.0 | 7.2574 | 52.0 | 1467.0 | 190.0 | |
| 3 | 341300.0 | 5.6431 | 52.0 | 1274.0 | 235.0 | |
| 4 | 342200.0 | 3.8462 | 52.0 | 1627.0 | 280.0 | |
| 5 | 269700.0 | 4.0368 | 52.0 | 919.0 | 213.0 | |
| 6 | 299200.0 | 3.6591 | 52.0 | 2535.0 | 489.0 | |
| 7 | 241400.0 | 3.1200 | 52.0 | 3104.0 | 687.0 | |
| 8 | 226700.0 | 2.0804 | 42.0 | 2555.0 | 665.0 | |
| 9 | 261100.0 | 3.6912 | 52.0 | 3549.0 | 707.0 | |

Dtypes, missing, duplicates

In [4]: `display(df.info())`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   median_house_value  20640 non-null   float64
 1   median_income       20640 non-null   float64
 2   housing_median_age  20640 non-null   float64
 3   total_rooms         20640 non-null   float64
 4   total_bedrooms      20640 non-null   float64
 5   population          20640 non-null   float64
 6   households          20640 non-null   float64
 7   latitude             20640 non-null   float64
 8   longitude            20640 non-null   float64
dtypes: float64(9)
memory usage: 1.4 MB
None
```

In [5]: `print("Dtypes, missing counts")`

```
dtypes_missing = pd.DataFrame({
    "dtype": df.dtypes.astype(str),
    "missing_count": df.isna().sum(),
    "missing_rate": (df.isna().mean()).round(4),
    "unique_count": df.nunique(dropna=True)
}).sort_values("missing_count", ascending=False)
display(dtypes_missing)

dups = df.duplicated().sum()
print("Duplicated rows:", dups)
```

Dtypes, missing counts

| | dtype | missing_count | missing_rate | unique_count |
|---------------------------|--------------|----------------------|---------------------|---------------------|
| median_house_value | float64 | 0 | 0.0 | 3842 |
| median_income | float64 | 0 | 0.0 | 12928 |
| housing_median_age | float64 | 0 | 0.0 | 52 |
| total_rooms | float64 | 0 | 0.0 | 5926 |
| total_bedrooms | float64 | 0 | 0.0 | 1928 |
| population | float64 | 0 | 0.0 | 3888 |
| households | float64 | 0 | 0.0 | 1815 |
| latitude | float64 | 0 | 0.0 | 862 |
| longitude | float64 | 0 | 0.0 | 844 |

Duplicated rows: 0

Statistics for numeric columns

```
In [6]: num = df.select_dtypes(include=[np.number])
if num.shape[1] > 0:
    display(num.describe().T)
else:
    print("No numeric columns found.")
```

| | count | mean | std | min | 25% |
|---------------------------|--------------|---------------|---------------|------------|-------------|
| median_house_value | 20640.0 | 206855.816909 | 115395.615874 | 14999.0000 | 119600.0000 |
| median_income | 20640.0 | 3.870671 | 1.899822 | 0.4999 | 2.5634 |
| housing_median_age | 20640.0 | 28.639486 | 12.585558 | 1.0000 | 18.0000 |
| total_rooms | 20640.0 | 2635.763081 | 2181.615252 | 2.0000 | 1447.7500 |
| total_bedrooms | 20640.0 | 537.898014 | 421.247906 | 1.0000 | 295.0000 |
| population | 20640.0 | 1425.476744 | 1132.462122 | 3.0000 | 787.0000 |
| households | 20640.0 | 499.539680 | 382.329753 | 1.0000 | 280.0000 |
| latitude | 20640.0 | 35.631861 | 2.135952 | 32.5400 | 33.9300 |
| longitude | 20640.0 | -119.569704 | 2.003532 | -124.3500 | -121.8000 |

```
In [7]: maxv = df['median_house_value'].max()
print("Max value:", maxv, " Count:", (df['median_house_value']==maxv).sum())
df['is_censored'] = (df['median_house_value']==maxv).astype(int)
df['is_censored'].value_counts()
```

Max value: 500001.0 Count: 965

```
Out[7]: is_censored
0    19675
1     965
Name: count, dtype: int64
```


```
In [8]: if num.shape[1] > 0:
    iqr = num.quantile(0.75) - num.quantile(0.25)
    range_ = num.max() - num.min()
    outlier_flag = (range_ / (iqr.replace(0, np.nan))).sort_values(ascending=False)
    print("\nTop features by range / IQR (may indicate heavy tails or outliers):")
    display(outlier_flag)
else:
    print("No numeric columns found.")
```

```
Top features by range / IQR (may indicate heavy tails or outliers):
population          38.037313
total_rooms         23.124835
households          18.710769
total_bedrooms      18.306818
median_income        6.651926
median_house_value   3.341960
housing_median_age   2.684211
longitude            2.649077
latitude             2.489418
dtype: float64
```

Correlation matrix for numeric features

```
In [9]: num = df.select_dtypes(include=[np.number]).drop(columns=['is_censored'], errors='ignore')

if num.shape[1] >= 2:
    corr = num.corr(numeric_only=True)
    plt.figure(figsize=(8,6))
    plt.imshow(corr, interpolation="nearest", aspect="auto")
    plt.colorbar()
    plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
    plt.yticks(range(len(corr.index)), corr.index)
    plt.title("Correlation matrix (numeric)")
    plt.tight_layout()
    plt.show()
else:
    print("Not enough numeric columns for correlation matrix.")
```



```
In [10]: target = "median_house_value"
if target in corr.columns:
    target_corr = corr[target].drop(target).abs().sort_values(ascending=False)
    print("Top features by absolute correlation with target:")
    display(target_corr.head(10))
else:
    print("Target not in numeric columns for correlation.")
```

Top features by absolute correlation with target:

| | |
|--------------------|----------|
| median_income | 0.688075 |
| latitude | 0.144160 |
| total_rooms | 0.134153 |
| housing_median_age | 0.105623 |
| households | 0.065843 |
| total_bedrooms | 0.050594 |
| longitude | 0.045967 |
| population | 0.024650 |

Name: median_house_value, dtype: float64

Target distribution


```
In [11]: plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
df[target].dropna().hist(bins=40)
plt.title(f"Histogram: {target}")
plt.xlabel(target)
```

```

plt.subplot(1,2,2)
df[target].dropna().plot.box(vert=False)
plt.title(f"Boxplot: {target}")

plt.tight_layout()
plt.show()

```



```

In [12]: skew = num.skew().sort_values(ascending=False)
display(skew)

```

| | |
|--------------------|-----------|
| population | 4.935858 |
| total_rooms | 4.147343 |
| total_bedrooms | 3.453073 |
| households | 3.410438 |
| median_income | 1.646657 |
| median_house_value | 0.977763 |
| latitude | 0.465953 |
| housing_median_age | 0.060331 |
| longitude | -0.297801 |
| dtype: | float64 |

```

In [13]: kurtosis = num.kurtosis().sort_values(ascending=False)
display(kurtosis)

```

| | |
|--------------------|-----------|
| population | 73.553116 |
| total_rooms | 32.630927 |
| households | 22.057988 |
| total_bedrooms | 21.923495 |
| median_income | 4.952524 |
| median_house_value | 0.327870 |
| housing_median_age | -0.800629 |
| latitude | -1.117760 |
| longitude | -1.330152 |
| dtype: | float64 |

Scatter and Pair plots for top numeric features correlated with target

```

In [14]: if target in num.columns:
    corr_with_target = num.corr()[target].abs().drop(target).sort_values(ascending=True)
    top_feats = corr_with_target.head(3).index.tolist()
    print("Top numeric features vs target:", top_feats)


```

```

n = len(top_feats)
if n > 0:
    plt.figure(figsize=(5*n,3))
    for i, f in enumerate(top_feats,1):
        plt.subplot(1,n,i)
        plt.scatter(df[f], df[target], s=8, alpha=0.6)
        plt.xlabel(f); plt.ylabel(target)
        plt.title(f"{f} vs {target}")
    plt.tight_layout()
    plt.show()
else:
    print("Target not numeric or no numeric features.")

```

Top numeric features vs target: ['median_income', 'latitude', 'total_rooms']




Region exploration

```
In [15]: coords = df[['longitude','latitude']]
km = KMeans(n_clusters=8, random_state=0).fit(coords)
df['region_id'] = km.labels_
df.groupby('region_id')['median_house_value'].median().sort_values()
```

```
Out[15]: region_id
2      76400.0
4      85800.0
6      89000.0
3     117800.0
7     169650.0
5     181600.0
0     198400.0
1     245100.0
Name: median_house_value, dtype: float64
```

```
In [16]: plt.figure(figsize=(8,6))
plt.scatter(df['longitude'], df['latitude'], c=df['region_id'], s=8, cmap='tab10',
plt.gca().invert_xaxis()
plt.title("KMeans clusters (region_id)")
plt.xlabel("longitude"); plt.ylabel("latitude")
plt.show()
```



```
In [17]: df.groupby('region_id', observed=False)[['median_house_value']].agg(['count','median'])
```

```
Out[17]:      count    median        mean
```


| region_id | | | |
|-----------|-------|----------|---------------|
| | count | median | mean |
| 2 | 1117 | 76400.0 | 86213.249776 |
| 4 | 685 | 85800.0 | 92102.916788 |
| 6 | 474 | 89000.0 | 106898.529536 |
| 3 | 2203 | 117800.0 | 127250.976396 |
| 7 | 1678 | 169650.0 | 198141.446365 |
| 5 | 986 | 181600.0 | 191438.676471 |
| 0 | 8652 | 198400.0 | 229496.027624 |
| 1 | 4845 | 245100.0 | 262594.436739 |

```
In [18]: centroids = km.cluster_centers_
plt.figure(figsize=(8,6))
plt.hexbin(df["longitude"], df["latitude"], gridsize=80, mincnt=1, cmap='Greys', alpha=0.5)
plt.scatter(centroids[:,1], centroids[:,0], c='red', s=100, marker='X')
for i,(lon,lat) in enumerate(zip(centroids[:,1], centroids[:,0])):
    plt.text(lon, lat, str(i))
```

```

plt.text(lon, lat, str(i), color='white', fontsize=9, ha='center', va='center')
plt.gca().invert_xaxis()
plt.show()

```



Categorical check

```

In [19]: categ = df.select_dtypes(include=["object", "category"]).columns.tolist()
if not categ:
    print("No categorical columns found.")
else:
    for c in categ:
        print(f"\nColumn: {c} | Unique: {df[c].nunique()}")
        print(df[c].value_counts().head(5))

```

No categorical columns found.

EDA zhrnutie:

- **Veľkosť a čistota dát:** 20 640 riadkov, 9 stĺpcov. V EDA žiadne chýbajúce hodnoty a 0 úplných duplicitných riadkov — dáta sú kompletne a pripravené na ďalšiu analýzu.
- **Ciel — median_house_value :**
 - mean ≈ **206 856**, median ≈ **179 700**, veľký rozptyl.
 - max = **500001.0**, počet pri maxime = **965**

- **Hlavné prediktory (lineárne):**

- `median_income` má najsilnejšiu pozitívnu koreláciu s cenou (~**0.69**) — kľúčový prediktor.
- Ďalej sú relevantné priestorové prvky (`latitude / longitude`) a štrukturálne metriky domu `total_rooms` ...

- **Outliery a šikmost'**:

- Stĺpce `population`, `total_rooms`, `households`, `total_bedrooms` majú veľký `range / IQR` a vysoký skew → silné chvosty.
- `median_house_value` skew ≈ **0.98** (pravostranný).

- **Vizualizácie a priestorový vzor:**

- Scatter / pairplot potvrdili silný trend `median_income` vs cena a regionálne pásy podľa `latitude`.
- Hexbin (longitude vs latitude) ukázal jasné **hotspoty**

- **Regionálne klastre (KMeans, k=8):**

- KMeans vytvoril regiony s rôznymi mediánmi
- Veľkosti klastrov sú nerovnomerné (niektoré stovky, iné tisícky riadkov) — pri použití v modeli zohľadni nerovnomernosť.
- Dôležité: KMeans v EDA možno fitnúť na celom df len pre exploráciu

Data preprocessing and normalization

```
In [20]: CONFIG={}
TARGET = "median_house_value"

EXCLUDE = {TARGET, "is_censored"}
if "region_id" in df.columns:
    EXCLUDE.add("region_id")

FEATURES_NUM = (
    df.drop(columns=[c for c in EXCLUDE if c in df.columns], errors="ignore")
        .select_dtypes(include=[np.number])
        .columns.tolist()
)

print(f"Target: {TARGET}")
print(f"Počet numerických features: {len(FEATURES_NUM)}")
print("Features:", FEATURES_NUM)
```

Target: median_house_value
Počet numerických features: 8
Features: ['median_income', 'housing_median_age', 'total_rooms', 'total_bedrooms',
'population', 'households', 'latitude', 'longitude']

```
In [21]: def suggest_log_cols(frame: pd.DataFrame, cols: list[str], skew_threshold: float =
    skew_vals = frame[cols].skew(numeric_only=True)
```

```

    return skew_vals[skew_vals.abs() >= skew_threshold].index.tolist()

def make_preprocess(feature_cols: list[str], log_cols: list[str] | None = None):
    if log_cols is None:
        log_cols = []
    other_cols = [c for c in feature_cols if c not in log_cols]

    col_tf = ColumnTransformer(
        transformers=[
            ("log", FunctionTransformer(np.log1p, validate=False), log_cols),
            ("num", "passthrough", other_cols),
        ],
        remainder="drop"
    )
    preprocess = Pipeline([
        ("cols", col_tf),
        ("scaler", StandardScaler()),
    ])
    return preprocess

```

In [22]: LOG_COLS_SUGGESTED = suggest_log_cols(df, FEATURES_NUM, skew_threshold=1.0)
print("Návrh log1p stĺpcov (na celej vzorke):", LOG_COLS_SUGGESTED)

Návrh log1p stĺpcov (na celej vzorke): ['median_income', 'total_rooms', 'total_bedrooms', 'population', 'households']

Data preprocessing & normalization

- **Leakage ochrana:** z modelu vylučujem `is_censored` (je odvodený z targetu) a dočasne aj `region_id` z EDA, aby sme netlačili informácie z celej vzorky do tréningu.
- **Výber numerických vstupov:** pracujem len s číselnými stĺpcami po odfiltrovaní `EXCLUDE`.
- **Normalizácia:** budeme používať `StandardScaler` (0-mean, jednotková odchýlka).
- **Stabilizácia rozdelení:** pre stĺpce s vysokou šikmostou budeme aplikovať `log1p`. Návrh šikmých stĺpcov je v tejto fáze informatívny — skutočné `log_cols` vyberieme **až na tréningovej sade** po splite, aby sme sa vyhli leakage.
- **Pipeline:** definoval som `make_preprocess(...)`, ktorá skombinuje (1) voliteľný `log1p` a (2) škálovanie, aby bol preproces konzistentný v train/val/test aj pri nasadení.

Data split

In [23]:

```

try:
    _ds = CONFIG.get("data_split", {})
    SEED = CONFIG.get("seed", 42)
    TEST_SIZE = _ds.get("test_size", 0.20)
    VAL_SIZE = _ds.get("val_size", 0.20)
    N_BINS = _ds.get("strat_bins", 10)
except NameError:
    SEED = 42
    TEST_SIZE = 0.20

```

```

VAL_SIZE = 0.20
N_BINS = 10

rng = np.random.RandomState(SEED)

```

In [24]:

```

X_all = df[FEATURES_NUM].copy()
y_all = df[TARGET].values

y_bins_all = pd.qcut(y_all, q=N_BINS, labels=False, duplicates="drop")

X_trval, X_test, y_trval, y_test, bins_trval, bins_test = train_test_split(
    X_all, y_all, y_bins_all,
    test_size=TEST_SIZE,
    random_state=SEED,
    stratify=y_bins_all
)

bins_trval_local = pd.qcut(y_trval, q=N_BINS, labels=False, duplicates="drop")

X_train, X_val, y_train, y_val = train_test_split(
    X_trval, y_trval,
    test_size=VAL_SIZE,
    random_state=SEED,
    stratify=bins_trval_local
)

print(f"Train: {X_train.shape}, Val: {X_val.shape}, Test: {X_test.shape}")

```

Train: (13209, 8), Val: (3303, 8), Test: (4128, 8)

In [25]:

```

def _brief_stats(name, y):
    return pd.Series({
        "n": len(y),
        "mean": float(np.mean(y)),
        "std": float(np.std(y)),
        "min": float(np.min(y)),
        "25%": float(np.quantile(y, 0.25)),
        "50%": float(np.quantile(y, 0.50)),
        "75%": float(np.quantile(y, 0.75)),
        "max": float(np.max(y)),
    }, name=name)

display(pd.concat([
    _brief_stats("train", y_train),
    _brief_stats("val", y_val),
    _brief_stats("test", y_test),
], axis=1))

```

| | train | val | test |
|-------------|---------------|---------------|---------------|
| n | 13209.000000 | 3303.000000 | 4128.000000 |
| mean | 206916.513286 | 207108.405086 | 206459.490068 |
| std | 115403.069340 | 115991.929765 | 114877.300401 |
| min | 14999.000000 | 22500.000000 | 14999.000000 |
| 25% | 119600.000000 | 118900.000000 | 119875.000000 |
| 50% | 179700.000000 | 179700.000000 | 179650.000000 |
| 75% | 264500.000000 | 266000.000000 | 264900.000000 |
| max | 500001.000000 | 500001.000000 | 500001.000000 |

Data split — čo a prečo

- **Ciel:** mať train/val/test s podobnou distribúciou ciela.
- **Metóda:** stratifikácia podľa kvantilov (`pd.qcut`) + `train_test_split` so `stratify` .
- **Pomer:** 20 % **test**, zvyšok rozdelený na **train/val** (20 % z trénovacej časti ide na validáciu).
- **Kontrola:** tabuľka štatistik (n, mean, kvantily) pre `y_train` , `y_val` , `y_test` .

Configuration

```
In [26]: ART_DIR = "artifacts"

Path(ART_DIR).mkdir(parents=True, exist_ok=True)

CONFIG = {
    "seed": SEED,
    "target": TARGET,
    "artifacts_dir": ART_DIR,
    "data_split": {
        "test_size": TEST_SIZE,
        "val_size": VAL_SIZE,
        "strat_bins": N_BINS
    },
    "preprocess": {
        "scaler": "standard",
        "auto_log_skew": True,
        "skew_threshold": 1.0,
        "manual_log_cols": []
    },
    "features": FEATURES_NUM
}
```

```
In [27]: if CONFIG["preprocess"]["auto_log_skew"]:
    log_cols_final = suggest_log_cols(
        X_train, CONFIG["features"], CONFIG["preprocess"]["skew_threshold"]
    )
else:
    log_cols_final = [c for c in CONFIG["preprocess"]["manual_log_cols"]
                      if c in CONFIG["features"]]

CONFIG["preprocess"]["log_cols"] = log_cols_final
print("log_cols (z TRAIN):", CONFIG["preprocess"]["log_cols"])

cfg_json = json.dumps(CONFIG, ensure_ascii=False, indent=2)
print(cfg_json)

log_cols (z TRAIN): ['median_income', 'total_rooms', 'total_bedrooms', 'population',
'households']
{
    "seed": 42,
    "target": "median_house_value",
    "artifacts_dir": "artifacts",
    "data_split": {
        "test_size": 0.2,
        "val_size": 0.2,
        "strat_bins": 10
    },
    "preprocess": {
        "scaler": "standard",
        "auto_log_skew": true,
        "skew_threshold": 1.0,
        "manual_log_cols": [],
        "log_cols": [
            "median_income",
            "total_rooms",
            "total_bedrooms",
            "population",
            "households"
        ]
    },
    "features": [
        "median_income",
        "housing_median_age",
        "total_rooms",
        "total_bedrooms",
        "population",
        "households",
        "latitude",
        "longitude"
    ]
}
```

Configuration — čo a prečo

- **Jedno miesto pravdy:** všetky dôležité nastavenia (seed, cesty, pomery splitu, preproces) sú v `CONFIG`.

- **Features:** explicitne definujem zoznam numerických vstupov po odfiltrovaní potenciálneho leakage.
- **Preprocessing parametre:** voľba škálovača a pravidlá pre `log1p`. Kľúčové je, že `log_cols vyberáme z TRAIN` podľa šiknosti ($|skew| \geq 1.0$), aby sme sa vyhli leakage.

Experiment tracking

```
In [28]: EXP_HISTORY = []

def log_experiment_mem(params: dict, metrics: dict, notes: str = ""):
    run = {
        "run_id": str(uuid.uuid4())[:8],
        "timestamp": int(time.time()),
        "seed": CONFIG["seed"],
        "notes": notes,
        **{f"p_{k}": v for k, v in params.items()},
        **{f"m_{k}": v for k, v in metrics.items()},
    }
    EXP_HISTORY.append(run)
    return run
```

```
In [29]: preprocess = make_preprocess(CONFIG["features"], CONFIG["preprocess"]["log_cols"])

Xtr = preprocess.fit_transform(X_train)
Xva = preprocess.transform(X_val)
Xte = preprocess.transform(X_test)

baseline = DummyRegressor(strategy="median")
baseline.fit(Xtr, y_train)

def reg_metrics(y_true, y_pred):
    rmse = float(np.sqrt(mean_squared_error(y_true, y_pred)))
    mae = float(mean_absolute_error(y_true, y_pred))
    return {"rmse": rmse, "mae": mae}

m_tr = reg_metrics(y_train, baseline.predict(Xtr))
m_va = reg_metrics(y_val, baseline.predict(Xva))
m_te = reg_metrics(y_test, baseline.predict(Xte))

display(pd.DataFrame({"train": m_tr, "val": m_va, "test": m_te}))
```

| | train | val | test |
|-------------|---------------|---------------|---------------|
| rmse | 118568.996827 | 119186.192321 | 117952.806055 |
| mae | 88317.944280 | 88749.307296 | 88153.726017 |

Experiment tracking

- Metriky (RMSE, MAE, pre train/val/test)

- Na konci zobrazím posledné záznamy z pamäte cez `DataFrame`

Experiments

```
In [30]: # Počítam, že tieto objekty už existujú z tvorho pipeline:
# CONFIG, Xtr, Xva, Xte, y_train, y_val, y_test, preprocess
for name in ["CONFIG","Xtr","Xva","Xte","y_train","y_val","y_test","preprocess"]:
    assert name in globals(), f"Chýba {name} - najprv spust časť s preprocess/split

# GPU only: ak nie je CUDA, zastavíme.
assert torch.cuda.is_available(), "CUDA GPU nie je dostupná. Spust notebook v GPU p
DEV = torch.device("cuda:0")

# rýchlosť/konzistencia
SEED = int(CONFIG.get("seed", 42))
random.seed(SEED); np.random.seed(SEED); torch.manual_seed(SEED); torch.cuda.manual_
torch.backends.cudnn.benchmark = True
try:
    torch.set_float32_matmul_precision("high") # PyTorch 2.x
except Exception:
    pass

print("GPU:", torch.cuda.get_device_name(0))
print("n_features:", Xtr.shape[1])

# fallback metriky, ak ich náhodou nemáš (inak použije tvoje)
if "reg_metrics" not in globals():
    def reg_metrics(y_true, y_pred):
        y_true = np.asarray(y_true); y_pred = np.asarray(y_pred)
        rmse = float(np.sqrt(np.mean((y_true - y_pred) ** 2)))
        mae = float(np.mean(np.abs(y_true - y_pred)))
        ss_res = np.sum((y_true - y_pred)**2)
        ss_tot = np.sum((y_true - np.mean(y_true))**2)
        r2 = float(1 - ss_res/(ss_tot + 1e-12))
        return {"rmse": rmse, "mae": mae, "r2": r2}
```

GPU: NVIDIA GeForce GTX 1650

n_features: 8

```
In [31]: class ResidualBlock(nn.Module):
    def __init__(self, in_dim, hidden_dim, batchnorm=True, dropout=0.0):
        super().__init__()
        self.bn1 = nn.BatchNorm1d(in_dim) if batchnorm else nn.Identity()
        self.fc1 = nn.Linear(in_dim, hidden_dim)
        self.act = nn.ReLU()
        self.bn2 = nn.BatchNorm1d(hidden_dim) if batchnorm else nn.Identity()
        self.drop = nn.Dropout(dropout) if dropout > 0 else nn.Identity()
        self.fc2 = nn.Linear(hidden_dim, in_dim)

    def forward(self, x):
        h = self.bn1(x)
        h = self.fc1(h); h = self.act(h)
        h = self.bn2(h); h = self.drop(h)
        h = self.fc2(h)
```

```

        return x + h

class TabularMLP(nn.Module):
    def __init__(self, in_dim, hidden_layers=(128,64,32),
                 batchnorm=True, dropout=0.0,
                 residual=False, bottleneck=False):
        super().__init__()

        dims = [in_dim] + list(hidden_layers)
        if bottleneck and len(hidden_layers) >= 2:
            down = list(hidden_layers)
            up = list(hidden_layers[:-1])[:-1]
            dims = [in_dim] + down + up

        feats = []
        for i in range(len(dims)-1):
            inp, out = dims[i], dims[i+1]
            feats += [nn.Linear(inp, out)]
            if batchnorm: feats += [nn.BatchNorm1d(out)]
            feats += [nn.ReLU()]
            if dropout > 0: feats += [nn.Dropout(dropout)]
            if residual and out >= 16:
                feats += [ResidualBlock(out, max(out//2, 16), batchnorm=batchnorm,
                                       self.backbone = nn.Sequential(*feats)
                                       self.head = nn.Linear(dims[-1], 1)

    def forward(self, x):
        return self.head(self.backbone(x)).squeeze(1)

```

```

In [32]: # === Dataloader + tréning/eval s mixed precision (AMP) na GPU ===
# OPRAVY:
# - nové AMP API: from torch.amp import autocast, GradScaler (žiadny deprec. warning)
# - SAFE metriky: ak preds/target obsahuje inf/NaN, vrátíme "zlý" výsledok, aby sa
#   nepadlo na chybu
from torch.amp import autocast, GradScaler

# zvolíme dtype pre AMP: keď GPU podporuje BF16, použijeme ho (stabilnejšie), inak
SUPPORTS_BF16 = getattr(torch.cuda, "is_bf16_supported", lambda: False)()
AMP_DTYPE = torch.bfloat16 if SUPPORTS_BF16 else torch.float16

def _to_tensor(x, y):
    return (torch.tensor(x, dtype=torch.float32),
            torch.tensor(y, dtype=torch.float32))

def make_loaders(Xtr, ytr, Xva, yva, batch_size=512):
    Xtr_t, ytr_t = _to_tensor(Xtr, ytr)
    Xva_t, yva_t = _to_tensor(Xva, yva)
    tr = TensorDataset(Xtr_t, ytr_t)
    va = TensorDataset(Xva_t, yva_t)
    return (DataLoader(tr, batch_size=batch_size, shuffle=True, pin_memory=True),
            DataLoader(va, batch_size=batch_size, shuffle=False, pin_memory=True))

def safe_reg_metrics(y_true, y_pred):
    # robustné metriky: ak sa tréning "rozletí", nepadneme na chybu
    yt = np.asarray(y_true, dtype=np.float64).ravel()
    yp = np.asarray(y_pred, dtype=np.float64).ravel()

```

```

        if not (np.all(np.isfinite(yt)) and np.all(np.isfinite(yp))):
            return {"rmse": 1e12, "mae": 1e12, "r2": -1.0}
        rmse = float(np.sqrt(np.mean((yt - yp) ** 2)))
        mae = float(np.mean(np.abs(yt - yp)))
        ss_res = float(np.sum((yt - yp) ** 2))
        ss_tot = float(np.sum((yt - np.mean(yt)) ** 2))
        r2 = float(1 - ss_res / (ss_tot + 1e-12))
        return {"rmse": rmse, "mae": mae, "r2": r2}

def _grad_scaler():
    # pri BF16 scaling netreba; pri FP16 škálujeme
    return GradScaler('cuda') if AMP_DTYPE == torch.float16 else None

def train_epoch(model, loader, opt, loss_fn, max_grad_norm=None, scaler=None):
    model.train(); losses=[]
    for xb, yb in loader:
        xb = xb.to(DEV, non_blocking=True); yb = yb.to(DEV, non_blocking=True)
        opt.zero_grad(set_to_none=True)
        with autocast('cuda', dtype=AMP_DTYPE):
            pred = model(xb)
            loss = loss_fn(pred, yb)
            if scaler is not None:
                scaler.scale(loss).backward()
                if max_grad_norm:
                    scaler.unscale_(opt)
                    nn.utils.clip_grad_norm_(model.parameters(), max_grad_norm)
                scaler.step(opt); scaler.update()
            else:
                loss.backward()
                if max_grad_norm:
                    nn.utils.clip_grad_norm_(model.parameters(), max_grad_norm)
                opt.step()
            losses.append(loss.item())
    return float(np.mean(losses))

@torch.no_grad()
def eval_epoch(model, loader, loss_fn):
    model.eval(); losses=[]; y_true=[]; y_pred=[]
    for xb, yb in loader:
        xb = xb.to(DEV, non_blocking=True); yb = yb.to(DEV, non_blocking=True)
        with autocast('cuda', dtype=AMP_DTYPE):
            pred = model(xb)
            loss = loss_fn(pred, yb)
            losses.append(loss.item())
            y_true.append(yb.cpu().numpy())
            y_pred.append(pred.float().cpu().numpy())
    y_true = np.concatenate(y_true); y_pred = np.concatenate(y_pred)
    return float(np.mean(losses)), safe_reg_metrics(y_true, y_pred), (y_true, y_pred)

DEFAULT_CFG = dict(
    seed=SEED,
    hidden_layers=[128, 64, 32],
    batchnorm=True,
    dropout=0.10,
    residual=False,
    bottleneck=True,
)

```

```

        optimizer="adam",    # "adam" / "rmsprop" / "sgd"
        lr=1e-3,
        weight_decay=1e-4,
        batch_size=512,
        epochs=200,
        patience=20,
        grad_clip=1.0
    )

def run_experiment(cfg, Xtr, ytr, Xva, yva, Xte, yte):
    torch.manual_seed(cfg["seed"]); torch.cuda.manual_seed_all(cfg["seed"])
    scaler = _grad_scaler()

    tr_loader, va_loader = make_loaders(Xtr, ytr, Xva, yva, batch_size=cfg["batch_size"])

    model = TabularMLP(
        in_dim=Xtr.shape[1],
        hidden_layers=tuple(cfg["hidden_layers"]),
        batchnorm=cfg["batchnorm"], dropout=cfg["dropout"],
        residual=cfg["residual"], bottleneck=cfg["bottleneck"]
    ).to(DEV)

    if cfg["optimizer"] == "adam":
        opt = torch.optim.Adam(model.parameters(), lr=cfg["lr"], weight_decay=cfg["weight_decay"])
    elif cfg["optimizer"] == "rmsprop":
        opt = torch.optim.RMSprop(model.parameters(), lr=cfg["lr"], weight_decay=cfg["weight_decay"])
    else:
        opt = torch.optim.SGD(model.parameters(), lr=cfg["lr"], momentum=0.9,
                              nesterov=True, weight_decay=cfg["weight_decay"])

    loss_fn = nn.MSELoss()
    sched = torch.optim.lr_scheduler.ReduceLROnPlateau(
        opt, mode="min", factor=0.5, patience=max(1, cfg["patience"]//3)
    )

    best_val = float("inf"); best_state=None; patience_left = cfg["patience"]; hist = []
    for epoch in range(1, cfg["epochs"]+1):
        tr_loss = train_epoch(model, tr_loader, opt, loss_fn, max_grad_norm=cfg["grad_clip"])
        va_loss, va_m, _ = eval_epoch(model, va_loader, loss_fn)
        sched.step(va_loss)

        history.append({"epoch": epoch, "train_loss": tr_loss, "val_loss": va_loss,
                        **{f"val_{k}": v for k,v in va_m.items()}})

        if va_loss < best_val - 1e-6:
            best_val, patience_left = va_loss, cfg["patience"]
            best_state = {k: v.detach().cpu().clone() if hasattr(v, "detach") else
                         {k:v} for k,v in model.state_dict().items()}
        else:
            patience_left -= 1
            if patience_left <= 0:
                break

    if best_state is not None:
        model.load_state_dict(best_state)

```

```

# full eval
def _full_eval(X, y):
    loader = DataLoader(TensorDataset(torch.tensor(X, dtype=torch.float32),
                                      torch.tensor(y, dtype=torch.float32)),
                        batch_size=cfg["batch_size"], shuffle=False, pin_memory=True)
    return eval_epoch(model, loader, loss_fn)

tr_loss, tr_m, _ = _full_eval(Xtr, ytr)
va_loss, va_m, _ = _full_eval(Xva, yva)
te_loss, te_m, (y_true_t, y_pred_t) = _full_eval(Xte, yte)

result = {"train": tr_m, "val": va_m, "test": te_m,
          "best_val_loss": best_val, "epochs_run": len(history)}
return model, history, result, (y_true_t, y_pred_t)

```

```

In [33]: # === Baseline MLP experiment (GPU) ===
BASE_CFG = DEFAULT_CFG | {
    "hidden_layers": [128, 64, 32],
    "batchnorm": True,
    "dropout": 0.10,
    "residual": False,
    "bottleneck": True,
    "optimizer": "adam",
    "lr": 1e-3,
    "weight_decay": 1e-4,
    "batch_size": 1024, # GPU zvládne väčší batch (skús aj 512/2048)
    "epochs": 200,
    "patience": 20,
    "grad_clip": 1.0,
}

model, hist, res, (y_true_t, y_pred_t) = run_experiment(
    BASE_CFG, Xtr, y_train, Xva, y_val, Xte, y_test
)

print(pd.DataFrame(res).round(4))

# voliteľné: log do tvojej internej histórie, ak ju máš
if "log_experiment_mem" in globals():
    params = {"model": "MLP-baseline", "n_features_after_pp": int(Xtr.shape[1]), **B}
    metrics = **{f"train_{k}": v for k,v in res["train"].items()},
              **{f"val_{k}": v for k,v in res["val"].items()},
              **{f"test_{k}": v for k,v in res["test"].items()}
    _ = log_experiment_mem(params, metrics, notes="MLP baseline (GPU)")

```

| | train | val | test | best_val_loss | epochs_run |
|------|-------------|-------------|-------------|---------------|------------|
| rmse | 236896.9037 | 237351.4010 | 236242.2631 | 5.666855e+10 | 100 |
| mae | 206897.9735 | 207089.6188 | 206441.2393 | 5.666855e+10 | 100 |
| r2 | -3.2139 | -3.1872 | -3.2291 | 5.666855e+10 | 100 |

```


In [34]: # (voliteľné) rýchle grafy k baseline - Pred vs True + Residuals
fig, ax = plt.subplots(1,2, figsize=(10,4))
ax[0].scatter(y_true_t, y_pred_t, s=8, alpha=0.5)
ax[0].set_title("Test: Predictions vs True"); ax[0].set_xlabel("True"); ax[0].set_y
resid = y_pred_t - y_true_t
ax[1].hist(resid, bins=50)

```

```

ax[1].set_title("Test residuals (Pred - True)")
plt.tight_layout(); plt.show()

```



```

In [ ]: # === Random Search (robust) ===
search_space = {
    "hidden_layers": [
        [256,128,64],
        [128,64,32],
        [256,128,64,32],
        [128,64,32,64,128],
    ],
    "dropout": [0.0, 0.05, 0.10, 0.20, 0.30],
    "batchnorm": [True, False],
    "residual": [False, True],
    "bottleneck": [False, True],
    "optimizer": ["adam", "rmsprop", "sgd"],
    "lr": [3e-4, 1e-3], # tip: drž LR konzervatívne, 3e-3 môže divergovať pri AMP
    "weight_decay": [0.0, 1e-5, 1e-4, 1e-3],
    "batch_size": [512, 1024, 2048],
    "patience": [10, 20, 30],
}

def _pick(k): return random.choice(search_space[k])

N_TRIALS = 12 # uprav podľa času
rows = []; best = None

for t in range(1, N_TRIALS+1):
    cfg = DEFAULT_CFG | {
        "hidden_layers": _pick("hidden_layers"),
        "dropout": _pick("dropout"),
        "batchnorm": _pick("batchnorm"),
        "residual": _pick("residual"),
        "bottleneck": _pick("bottleneck"),
        "optimizer": _pick("optimizer"),
        "lr": _pick("lr"),
        "weight_decay": _pick("weight_decay"),
        "batch_size": _pick("batch_size"),
        "patience": _pick("patience"),
        "seed": random.randint(1, 10_000),
    }

```

```

try:
    model_i, hist_i, res_i, _ = run_experiment(cfg, Xtr, y_train, Xva, y_val, Xte)
    row = {"trial": t,
            "val_rmse": res_i["val"]["rmse"], "val_mae": res_i["val"]["mae"], "val_r2": res_i["val"]["r2"],
            "test_rmse": res_i["test"]["rmse"], "test_mae": res_i["test"]["mae"], "test_r2": res_i["test"]["r2"],
            "cfg": cfg, "error": ""}
except Exception as e:
    # ak sa aj tak niečo pokazi (napr. numerické NaN), zaznač trial ako neúspešný
    row = {"trial": t, "val_rmse": 1e12, "val_mae": 1e12, "val_r2": -1.0,
            "test_rmse": 1e12, "test_mae": 1e12, "test_r2": -1.0, "cfg": cfg, "error": ""}

rows.append(row)

if best is None or row["val_rmse"] < best["val_rmse"]:
    best = row

res_df = pd.DataFrame(rows).sort_values("val_rmse").reset_index(drop=True)
display(res_df[["trial", "val_rmse", "val_mae", "val_r2", "test_rmse", "test_mae", "test_r2"]])

print("\nNajlepší trial (podľa val_rmse):")
display(res_df.iloc[0])

```

```

In [ ]: report = [{"model": "MLP-baseline", **{f"train_{k}": v for k, v in res["train"].items()},
                  **{f"val_{k}": v for k, v in res["val"].items()},
                  **{f"test_{k}": v for k, v in res["test"].items()}}

if "res_df" in globals() and len(res_df) > 0:
    top = res_df.iloc[0]
    report.append({"model": "MLP-best-random",
                   "train_rmse": np.nan, "train_mae": np.nan, "train_r2": np.nan,
                   "val_rmse": float(top["val_rmse"]), "val_mae": float(top["val_mae"]),
                   "test_rmse": float(top["test_rmse"]), "test_mae": float(top["test_mae"])})

report_df = pd.DataFrame(report)
display(report_df.round(4))

md = f"""# Houses - MLP (GPU) - Report

**Techniky:** BatchNorm, Dropout, Residual/Skip, Bottleneck, EarlyStopping, ReduceLROnPlateau

**Baseline config:**
```json
{json.dumps(BASE_CFG, indent=2)}
```

```

```

In [ ]: fig, ax = plt.subplots(1, 2, figsize=(10, 4))
ax[0].scatter(y_true_t, y_pred_t, s=8, alpha=0.5)
ax[0].set_title("Test: Predictions vs True"); ax[0].set_xlabel("True"); ax[0].set_ylabel("Pred")
resid = y_pred_t - y_true_t
ax[1].hist(resid, bins=50)
ax[1].set_title("Test residuals (Pred - True)")
plt.tight_layout(); plt.show()

```

In []:

